O'REILLY LIVE TRAINING · NOVEMBER 2026

Getting Started with *LangChain*

Build AI agents with LangChain 1.0 — from a hand-rolled loop to a deployed app.

Lucas Soares · Automata Learning Lab

HOW THIS WORKS

Training rhythm

Presentation
Concepts, short

→

Notebook demo
Live code

→

Q&A + summary
Polls & questions

→

Exercise
Optional, during Q&A

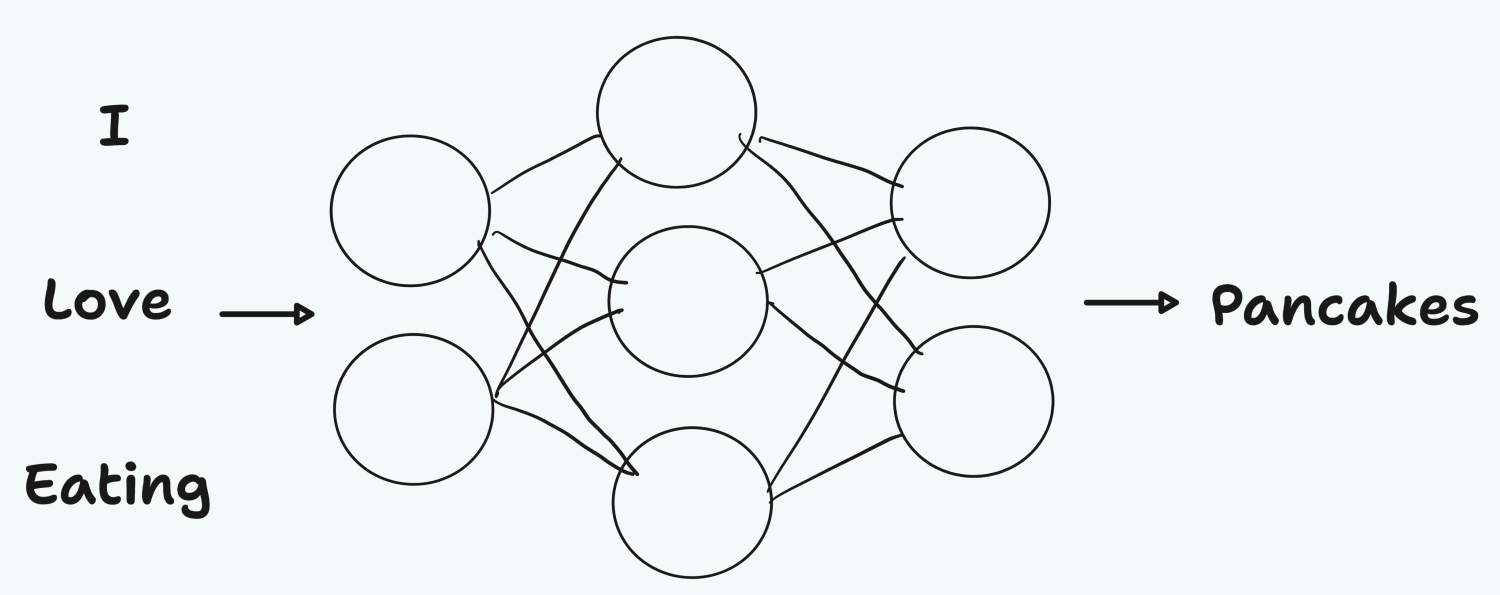
- Repeat for each part — interrupt with questions anytime

AGENDA

What we'll cover

1. **What is an agent?** — build one from scratch, then with `create_agent()` · *notebook 00*
2. **LangChain fundamentals** — models, messages, tools, memory, streaming · *notebook 1.0*
3. **Structured outputs** — typed results with Pydantic · *notebook 2.0*
4. **RAG** — embeddings, vector stores, retrieval agents · *notebook 3.0*
5. **Ship it** — `langgraph dev`, a fully local agent, and a public Vercel app · *demos*

Large Language Models predict the next word



What is LangChain?

- **"LangChain is the agent framework; LangGraph is the orchestration runtime"** — docs.langchain.com
- LangChain **1.x**: the agent is the central abstraction — `create_agent()`
- Built on **LangGraph** underneath — persistence, streaming, and deployment for free
- Legacy patterns (`AgentExecutor`, `prompt | llm | parser` piping) moved out to `langchain-classic` — not the story anymore

PART 0 · NOTEBOOK 00

What is an *agent*?

Build one by hand first — so the abstraction is a shortcut, not a black box.

"An agent is a model calling tools in a loop until a given task is complete."

— the official LangChain docs · "Agent = Model + Harness"

The core components of an agent

LLM

The reasoning engine — decides what to say, or which tool to call.

Tools

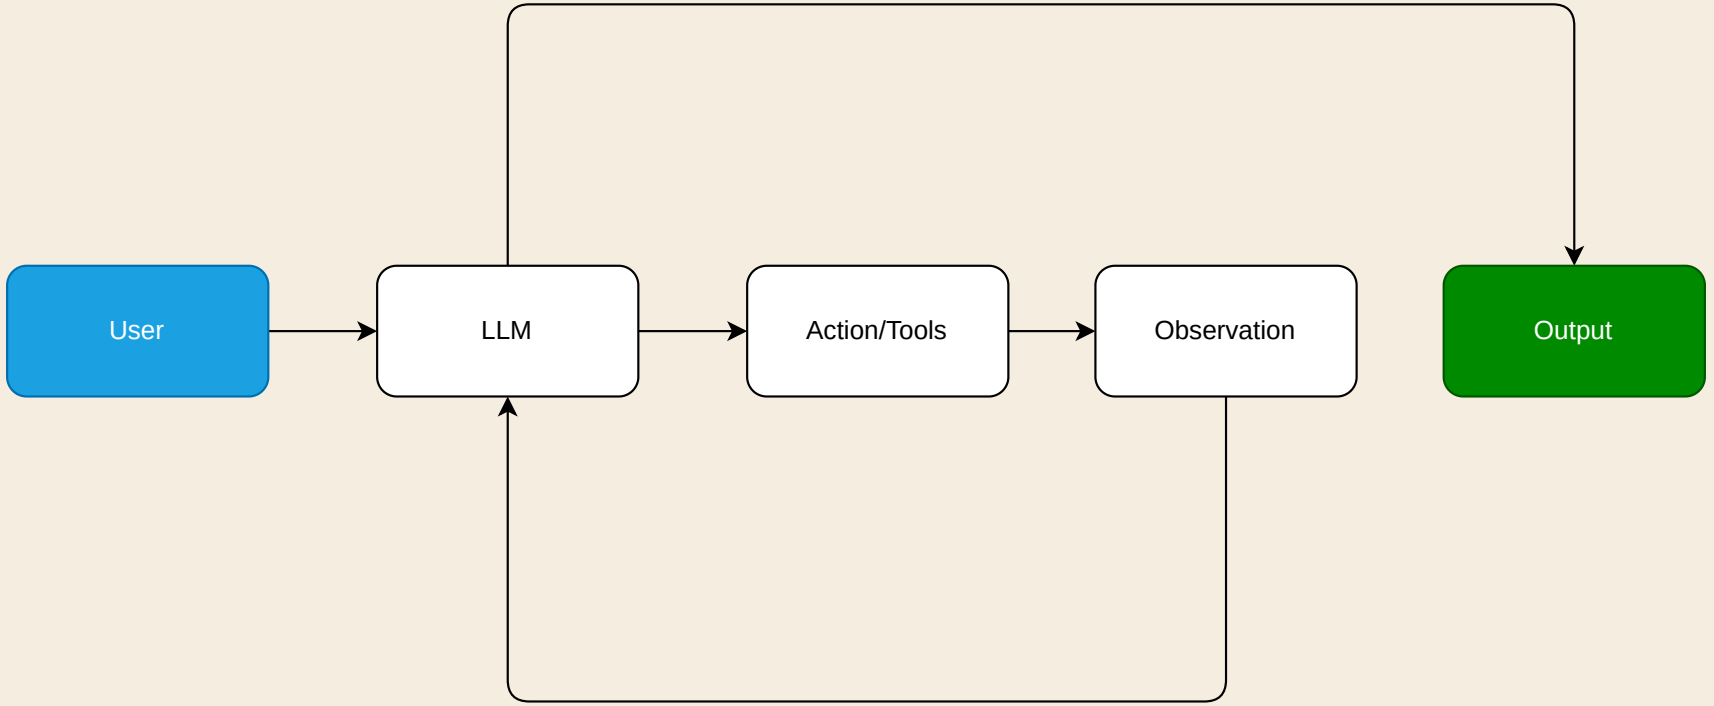
Plain Python functions. The LLM asks; *your code* executes.

Loop

Call model → run requested tools → feed results back → repeat.

- Every agent framework is a packaging of these three pieces

The agent loop



STEP 1 — JUST AN LLM CALL

```python
from langchain.chat_models import init_chat_model
from langchain.messages import HumanMessage, SystemMessage

model = init_chat_model("openai:gpt-6-luna", use_responses_api=True)

response = model.invoke([
    SystemMessage(content="You tell jokes."),
    HumanMessage(content="Tell me a joke about a bald teacher."),
])
print(response.text)   # an AIMessage — read its text with .text
```

No tools, no loop — the model can only *talk*. Let's give it hands.

STEP 2 — TOOLS WITH @TOOL

```python
from langchain_core.tools import tool

@tool
def create_folder(folder_name: str) → str:
    """Create a folder (directory) inside the sandbox workspace."""
    path = os.path.join(SANDBOX_DIR, folder_name)
    os.makedirs(path, exist_ok=True)
    return f"Folder created at: {path}"

tools = [create_folder, create_file, read_file, calculator]
tools_by_name = {t.name: t for t in tools}
```

Type hints + docstring **are** the schema the model sees.

STEP 3 — THE MODEL ASKS, YOUR CODE EXECUTES

```
model_with_tools = model.bind_tools(tools)

ai_message = model_with_tools.invoke(messages)
ai_message.tool_calls
# [{"name": "create_folder", "args": {"folder_name": "notes"}, "id": "..."}]

from langchain.messages import ToolMessage
messages.append(ai_message)
for call in ai_message.tool_calls:
    result = tools_by_name[call["name"]].invoke(call["args"])
    messages.append(ToolMessage(content=str(result), tool_call_id=call["id"]))
```

The model only *requested* the call — nothing ran until our loop ran it.

STEP 4 — THE WHOLE AGENT, BY HAND

```python
def run_agent_loop(user_task: str, max_iters: int = 5):
    messages = [SystemMessage(...), HumanMessage(content=user_task)]
    for i in range(max_iters):
        ai_message = model_with_tools.invoke(messages)
        messages.append(ai_message)
        if not ai_message.tool_calls:          # done — no more tool requests
            return ai_message.text
        for call in ai_message.tool_calls:     # execute each requested tool
            result = tools_by_name[call["name"]].invoke(call["args"])
            messages.append(ToolMessage(content=str(result),
                                        tool_call_id=call["id"]))
    return messages[-1].content
```

~15 lines. This **is** an agent.

STEP 5 — THE SAME AGENT WITH CREATE_AGENT()

```
from langchain.agents import create_agent

agent = create_agent(
    model=model,          # binds the tools for you
    tools=tools,
    system_prompt="You are a desktop file-system assistant.",
)

result = agent.invoke({"messages": "Create a folder called 'trip-notes' ..."})
result["messages"][-1].pretty_print()
```

Same behavior, three lines — now you know exactly what it's hiding.

What `create_agent()` replaced

From-scratch piece	`create_agent()`
`model.bind_tools(tools)`	internal, via `tools=[...]`
Building `System` / `Human` messages	`system_prompt=` + `{"messages": ...}`
Reading `tool_calls`, dispatching	internal
Building `ToolMessage(...)`	internal
The loop + safety cap	built into the compiled graph

NOTEBOOK DEMO

`00-langchain-basics-intro.ipynb`

Agents from scratch → `create_agent()`

Q&A · Part 0 recap

- An **agent = LLM + tools + loop** — the LLM requests, your code executes, results feed back
- `bind_tools()` makes the model emit `tool_calls`; `ToolMessage` carries results back
- `create_agent()` **packages that exact loop** — an abstraction you've now built yourself

Optional exercise

Add a `delete_file` tool to the from-scratch loop and have the agent clean up its own sandbox.

Poll

```
Which best describes what create_agent() returns in LangChain 1.0?

  A. A REST API endpoint for the model
  B. A compiled LangGraph graph you invoke with {"messages": [...]}
  C. A fixed script that cannot be modified once generated
  D. A raw string containing the model's response
```

PART 1 · NOTEBOOK 1.0

LangChain *fundamentals*

Six building blocks you'll use in every app.

The six building blocks

Agents

`create_agent()` — the central API

Models

`init_chat_model()`

Messages

system · human · ai · tool

Tools

`@tool` decorator

Memory

checkpointer + thread

Streaming

values · messages · custom

Models — one factory, any provider

Provider	Model string	Package
OpenAI	`openai:gpt-6-luna`	`langchain-openai`
Anthropic	`anthropic:claude-haiku-5-5`	`langchain-anthropic`
Google	`google_genai:gemini-3.5-flash-lite`	`langchain-google-genai`
Local (Ollama)	`ollama:gemma4`	`langchain-ollama`

- Three call styles: **`invoke()`** one response · **`stream()`** token by token · **`batch()`** parallel

MESSAGES — TYPED CONVERSATION TURNS

```python
from langchain_core.messages import SystemMessage, HumanMessage

messages = [
    SystemMessage("You are a pirate that answers briefly."),
    HumanMessage("What is machine learning?"),
]
response = model.invoke(messages)      # → AIMessage

response.usage_metadata               # token accounting
response.content_blocks               # v1: unified content across providers
```

Dict form works too: `{"role": "user", "content": "..."}` — fully interchangeable.

Tools — two hard rules

- **Type hints are mandatory** — they define the argument schema
- **Docstrings are mandatory** — they tell the model *when* to call it
- Complex inputs? Attach a Pydantic model: `@tool(args_schema=WeatherQuery)`
- Runtime context without hardcoding: `context_schema=` + `ToolRuntime` (dependency injection)

SHORT-TERM MEMORY — CHECKPOINTER + THREAD_ID

```
from langgraph.checkpoint.memory import InMemorySaver

agent = create_agent(
    model=model, tools=[search],
    checkpointer=InMemorySaver(),
)

config = {"configurable": {"thread_id": "session-1"}}
agent.invoke({"messages": "Hi! My name is Lucas."}, config)
agent.invoke({"messages": "What's my name?"}, config)   # remembers
# thread_id "session-2" → a fresh conversation, no Lucas
```

Streaming — three modes

Mode	You get	Use for
`"values"`	full state after each step	inspecting agent steps
`"messages"`	`(token, metadata)` tuples	chat UX, token by token
`"custom"`	your own events via `get_stream_writer()`	tool progress updates

"This course uses OpenAI's Responses API: `use_responses_api=True`, the API OpenAI recommends for new projects."

— the course gotcha: `ChatOpenAI` defaults to Chat Completions, where gpt-6-luna only calls tools with `reasoning_effort="none"` and some tiers can't call tools at all.
And `.content` is now a list of blocks, so read output with `.text`

NOTEBOOK DEMO

`1.0-langchain-fundamentals.ipynb`

Models · messages · tools · memory · streaming

Q&A · Part 1 recap

- `init_chat_model("provider:model")` — same call shape for OpenAI, Anthropic, Google, Ollama
- **Tools** are functions with type hints + docstrings; schemas are generated for you
- **Memory** = `checkpointer` + `thread_id`; **streaming** = `values` / `messages` / `custom`

Optional exercise

Create an agent with `create_agent()` and a tool that builds a schedule from a `task | date` table.

Poll

Which statement best captures the primary advantage of using a framework like LangChain for LLM-based applications?

A. It automatically trains your model without any configuration
B. It offers unified abstractions (agents, models, tools) that streamline development
C. It focuses solely on deployment, not on application logic
D. It restricts user inputs to enhance security

PART 2 · NOTEBOOK 2.0

Structured *outputs*

From "nice prose" to data your code can use.

Why structured output?

Free text

- "Maria seems like a strong candidate, with solid Python..."
- Can't sort, filter, or store it
- Parsing prose = regex pain

Structured

- `evaluation.fit_score` → `87`
- `evaluation.recommendation` → `"hire"`
- Type-safe: Pydantic validates every field

RESPONSE_FORMAT + PYDANTIC

```python
from pydantic import BaseModel, Field
from typing import List, Literal

class MovieReview(BaseModel):
    """Structured analysis of a movie review."""
    title: str
    sentiment: Literal["positive", "negative", "mixed"]
    rating: int = Field(ge=1, le=10)
    key_points: List[str]
    recommended: bool

model = init_chat_model("openai:gpt-6-luna", use_responses_api=True)
agent = create_agent(model=model, response_format=MovieReview)
result = agent.invoke({"messages": "Analyze this review: ..."})
review = result["structured_response"]     # a validated MovieReview
```

Two strategies under the hood

	`ProviderStrategy`	`ToolStrategy`
How	provider's native structured-output API	tool-calling to extract
Works with	OpenAI · Anthropic · Gemini	any tool-capable model
Speed	fastest	one extra hop
Gotcha	—	needs tool calling: `use_responses_api=True` for gpt-6-luna

- Pass the schema directly (`response_format=MovieReview`) and LangChain **auto-picks**

Real application: resume analyzer

- **Nested models**: `JobFitEvaluation` contains a `CandidateProfile` + a list of `SkillAssessment`
- Constrained fields: `fit_score: int = Field(ge=0, le=100)`, `recommendation: Literal["strong_hire", "hire", "maybe", "no_hire"]`
- **Batch scoring**: evaluate N resumes, then `sorted(evaluations, key=lambda e: e.fit_score)` — *only possible with structured output*

NOTEBOOK DEMO

`2.0-structured-outputs.ipynb`

Typed results, nested models, batch scoring

Q&A · Part 2 recap

- `response_format=YourModel` on `create_agent()` → `result["structured_response"]`
- **Pydantic** for production (validation, nesting); TypedDict for quick dicts
- `ProviderStrategy` vs `ToolStrategy` — auto-selection picks for you

Optional exercise

Define a Pydantic schema for meeting minutes (attendees, decisions, action items) and extract it from a transcript.

Poll

With response_format=MyPydanticModel, where does the typed result appear?

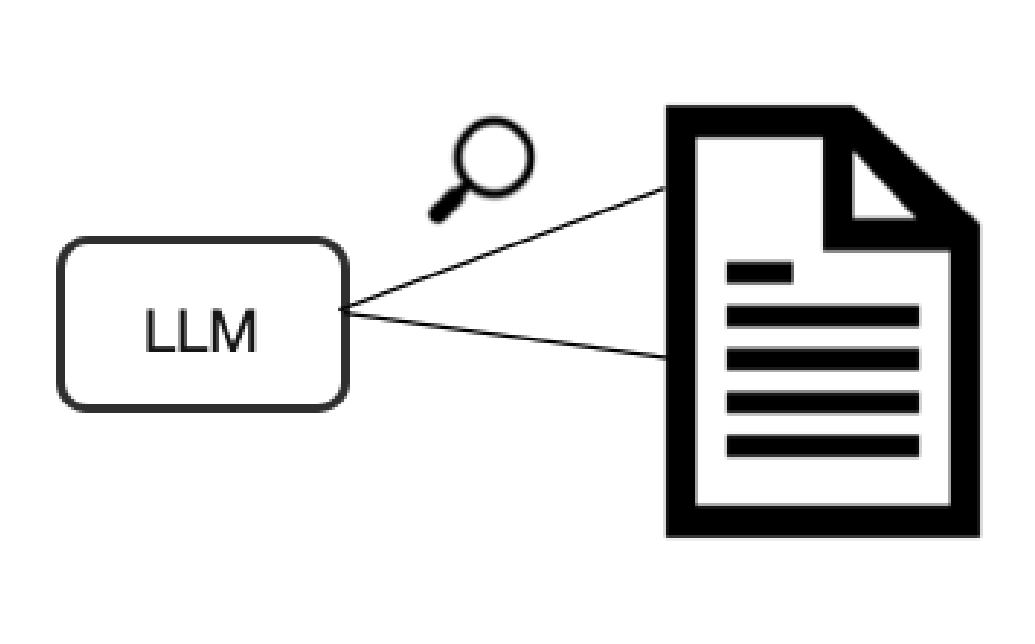
A. result["messages"][-1].text, as a JSON string to parse
B. result["structured_response"], as a validated model instance
C. It's written to a file next to the notebook
D. In a separate .schema attribute on the agent

PART 3 · NOTEBOOK 3.0

Chat over documents — *RAG*

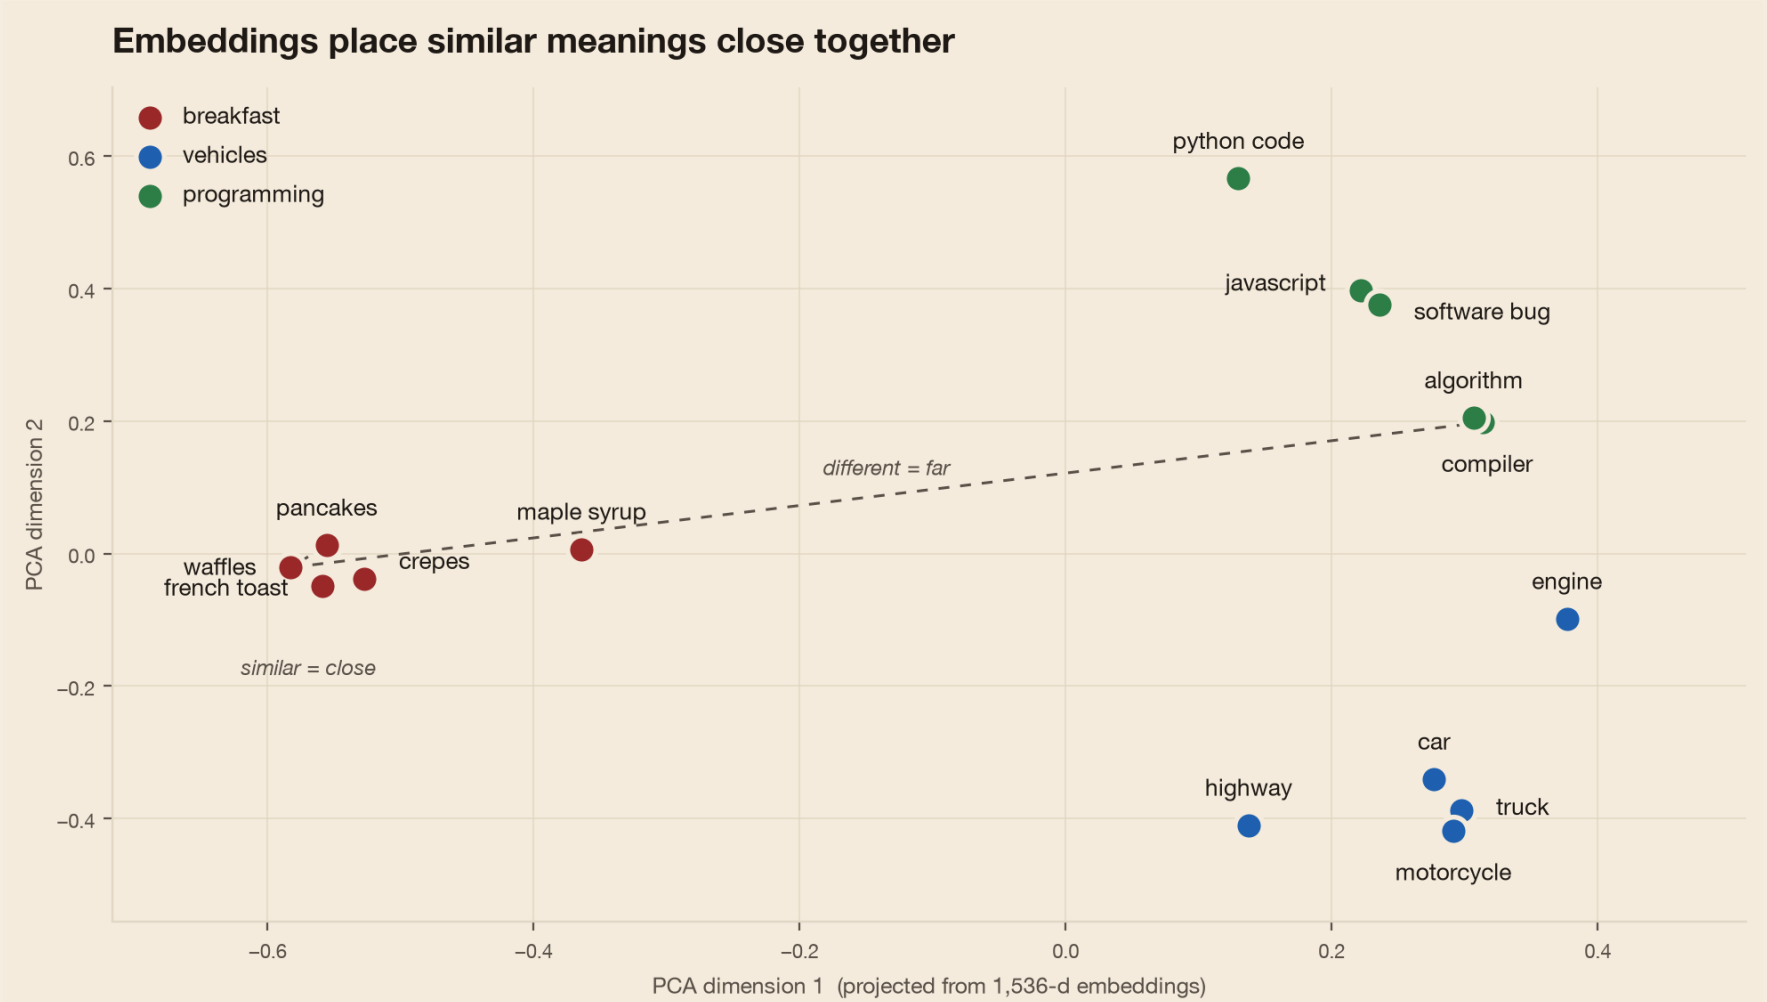
Retrieval Augmented Generation with LangChain.

Why RAG?



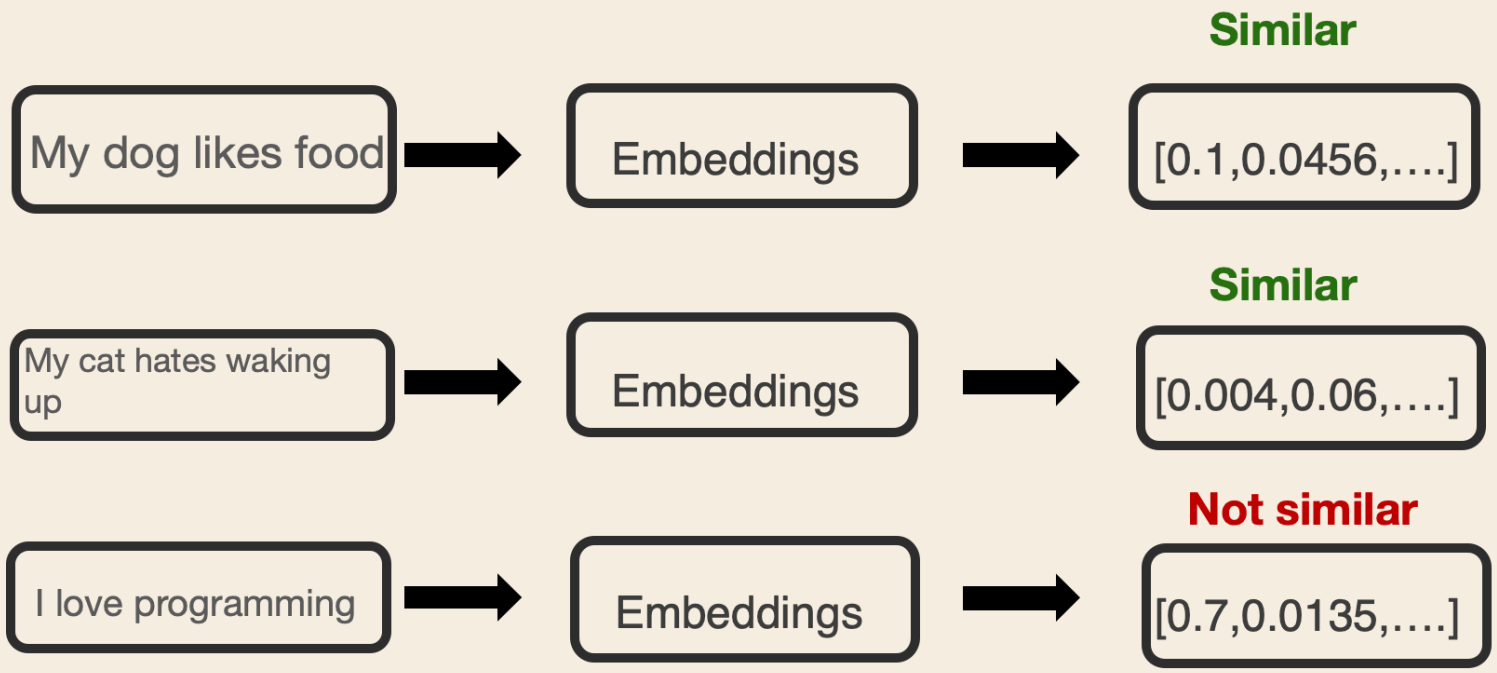
- Connect LLMs to **your** documents — PDFs, HTML, text
- Context windows are finite; your document base isn't
- The workaround: **embeddings** + retrieval of just-the-relevant chunks

Embeddings — meaning as geometry



- Real 1,536-d vectors from `text-embedding-3-small`, projected to 2D — distance ≈ meaning

Embeddings power similarity search



- A **vector store** (Chroma) indexes the vectors so "nearest neighbors" is one query

Indexing pipeline

Load
WebBaseLoader

→

Split
1000 chars, 200 overlap

→

Embed
text-embedding-3-small

→

Store
Chroma

```
splits = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=200).split_documents(docs)
vector_store = Chroma.from_documents(documents=splits,
    embedding=OpenAIEmbeddings(model="text-embedding-3-small"),
    collection_name="rag-tutorial", persist_directory="./chroma-db")
```

A RAG AGENT — RETRIEVAL AS A TOOL

```python
@tool(response_format="content_and_artifact")
def retrieve_context(query: str):
    """Retrieve information to help answer a query."""
    docs = vector_store.similarity_search(query, k=2)
    serialized = "\n\n".join(
        f"Source: {d.metadata}\nContent: {d.page_content}" for d in docs)
    return serialized, docs

agent = create_agent(model, [retrieve_context],
    system_prompt="Use the retrieval tool to answer user queries.")
```

The **agent decides** when (and how often) to search.

RAG agents vs RAG chains

	RAG agent — "Agentic RAG"	RAG chain — "2-Step RAG"
Search happens	when the LLM decides	always, before the call
Multiple searches	yes, with rewritten queries	one
Latency / cost	≥ 2 model calls	1 model call
How	retrieval `@tool`	`@dynamic_prompt` middleware
Best for	complex, multi-hop questions	predictable, speed-critical

- **Middleware** is v1's headline extension point — here it injects context and can **return source documents** for citations (more built-ins: summarization, human-in-the-loop, PII)

NOTEBOOK DEMO

`3.0-rag-with-langchain.ipynb`

Index two engineering blog posts, then chat with them

Q&A · Part 3 recap

- **Embeddings** turn text into vectors where distance ≈ similarity of meaning
- **Indexing**: load → split → embed → store (Chroma + `text-embedding-3-small`)
- **RAG agent** (LLM decides, retrieval `@tool`) vs **RAG chain** (always retrieve, middleware)

Optional exercise

Build a RAG agent with `create_agent()` over a PDF or CSV of your choice.

Poll

What is the primary purpose of embeddings when building LLM-powered document chat systems?

A. Managing API requests and rate limits
B. Turning text into numerical vectors for similarity comparison
C. Automatically generating user interface elements
D. Serving as a secure authentication method

PART 4 · DEMOS

From notebook to *app*

Three ways to ship the same agent.

1 · `langgraph dev` + Agent Chat UI

- The `demo/` app: RAG over two blog posts + Tavily web search — one `create_agent()` call
- `uv run langgraph dev` → starts the **Agent Server** at `http://localhost:2024`, hot-reloading, **persistence built in**
- Connect agentchat.vercel.app → graph ID `chat_over_docs`
- `LANGSMITH_TRACING=true` → every run traced in LangSmith (tools, tokens, full history)

2 · A fully local agent — `local_agent.py`

- `init_chat_model("ollama:gemma4")` — **no API key**, same `create_agent()` API
- File tools + `bash`, all **sandboxed**: path-escape guard, command blocklist, 20s timeout
- Tavily search + a running `memory.md` log the agent writes after each task
- One file, `uv run` — inline script dependencies

```
llm = init_chat_model("ollama:gemma4")
tools = [search, read_file, write_file, edit_file, delete_file, bash]
agent = create_agent(model=llm, tools=tools)
```

3 · Chat-over-PDFs on Vercel

- Same agent, now a **public web app**: chat pane + PDF viewer that jumps to the cited page
- Serverless can't host Chroma → index **pre-built offline** to JSON; runtime = numpy cosine similarity
- One-file FastAPI app: agent + `/api/chat` + embedded UI · full recipe in `internal-docs-search-agent/DEPLOYMENT_GUIDE.md`
- Torn down between classes — the guide recreates it in minutes

LIVE DEMO

`langgraph dev` · local agent · Vercel

The same `create_agent()` pattern, three deployment stories

Final Q&A · the whole course on one slide

- **Agent = LLM + tools + loop** — you built it by hand, then let `create_agent()` package it
- **Six building blocks**: agents, models, messages, tools, memory, streaming
- **Structured output** makes agent results programmable; **RAG** grounds them in your documents
- The same agent runs in a notebook, behind `langgraph dev`, fully local on Ollama, or on Vercel

THANK YOU

Thanks for *building* along

Questions → lucas@automata-learning-lab · course repo has every notebook, demo, and deploy guide

REFERENCES

Keep going

- LangChain documentation
- LangChain agents guide
- LangGraph · LangSmith
- Agent Chat UI
- ReAct paper · Toolformer
- Karpathy on agents
- Course repo: notebooks `00` → `3.0`, `demo/`, `local_agent.py`, `internal-docs-search-agent/`